The weather report tells you how frequently your build is broken or passing. By itself this data is interesting because you ultimately want passing builds so you can get your code into production and in front of your consumers. But it's not so bad to break builds from time to time because that could indicate that your team is taking risks—and calculated risks are good a lot of the time. You definitely don't want unhealthy or stormy builds because that could indicate a host of problems. As with CLOC from SCM, you should combine this with the other data you've collected so far to help get to the root of problems and understand why you're seeing the trends that you are.

You can get a lot more data from your CI system if you're generating reports from the different parts of your build process and publishing results. Earlier we talked about generating reports from your build script; this is where those reports become accessible and useful in the context of CI.

5.4. Key CI metrics: spotting trends in your data

The most basic information you can get from CI is successful and failed builds. Even though these seem like straightforward metrics, there's some ambiguity in interpreting them.

The most obvious problem is if your build fails all or most of the time; if that happens there's obviously a big problem. Conversely, if your build passes all the time, that might mean you have an awesome team, but it could also mean any of the following:

- Your team isn't doing any meaningful work.
- You don't have any tests that are running.
- Quality checks are disabled for your build.

When you run up against this issue, it's not a bad idea to pull a bit more data to sanity-check your build history. In this case you can take a look at the details from your CI builds to explore some of the following:

- Test reports