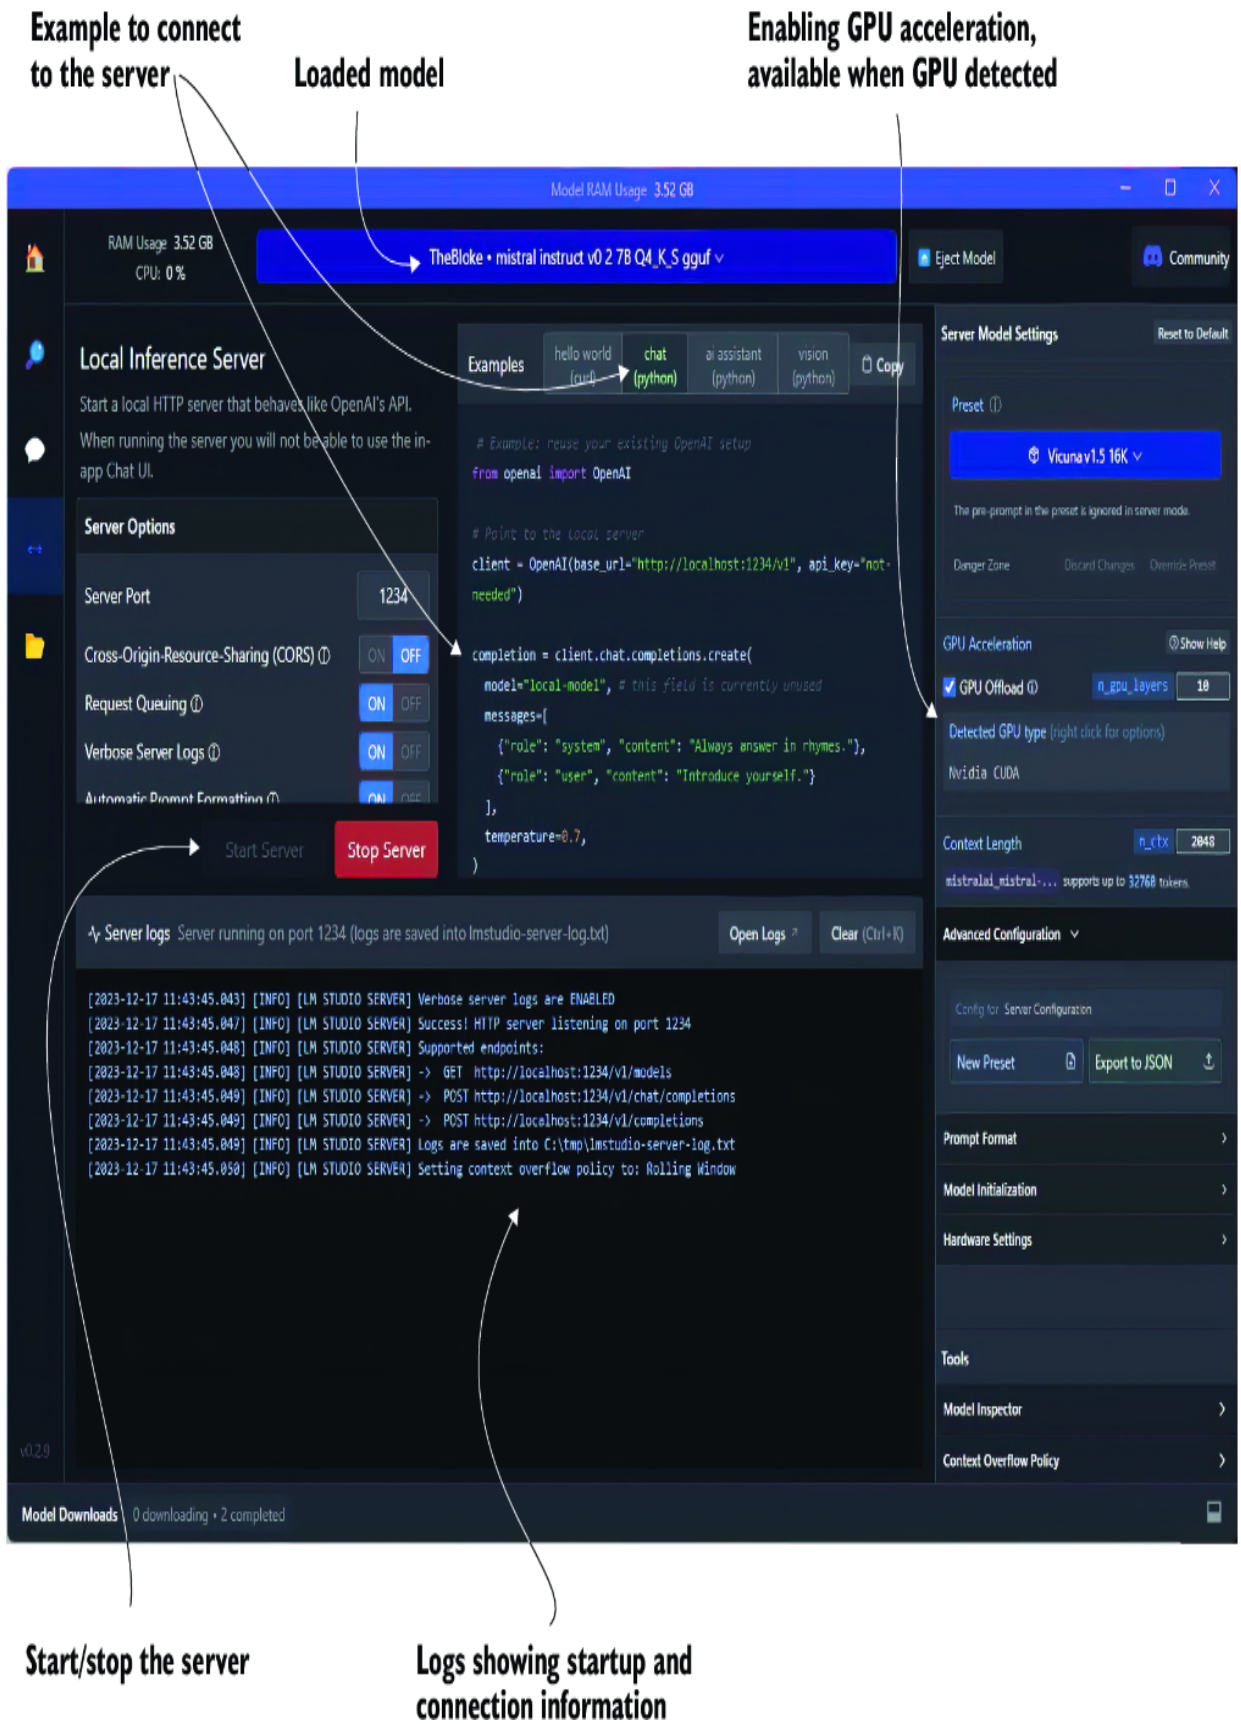

Example to connect to the server
Loaded model
Enabling GPU acceleration, available when GPU detected
Model RAM Usage 3.52 GB
RAM Usage 3.52 GB
CPU: 0 %
TheBloke • mistral instruct v0 2 7B Q4_K_S gguf
Eject Model
Community
Local Inference Server
Start a local HTTP server that behaves like OpenAI's API.
When running the server you will not be able to use the in-app Chat UI.
Server Options
Server Port
1234
Cross-Origin-Resource-Sharing (CORS)
Request Queuing
Verbose Server Logs
Start Server
Stop Server
Examples
hello world (curl)
chat (python)
ai assistant (python)
vision (python)
Copy
Example: reuse your existing OpenAI setup
from openai import OpenAI
Point to the local server
client = OpenAI(base_url="http://localhost:1234/v1", api_key="not-needed")
completion = client.chat.completions.create(
model="local-model", # this field is currently unused
messages=[
{"role": "system", "content": "Always answer in rhymes."},
{"role": "user", "content": "Introduce yourself."}
],
temperature=0.7,
)
Server logs Server running on port 1234 (logs are saved into lmstudio-server-log.txt)
Open Logs
Clear (Ctrl+K)
[2023-12-17 11:43:45.043] [INFO] [LM STUDIO SERVER] Verbose server logs are ENABLED
[2023-12-17 11:43:45.047] [INFO] [LM STUDIO SERVER] Success! HTTP server listening on port 1234
[2023-12-17 11:43:45.048] [INFO] [LM STUDIO SERVER] Supported endpoints:
[2023-12-17 11:43:45.048] [INFO] [LM STUDIO SERVER] -> GET http://localhost:1234/v1/models
[2023-12-17 11:43:45.049] [INFO] [LM STUDIO SERVER] -> POST http://localhost:1234/v1/chat/completions
[2023-12-17 11:43:45.049] [INFO] [LM STUDIO SERVER] -> POST http://localhost:1234/v1/completions
[2023-12-17 11:43:45.049] [INFO] [LM STUDIO SERVER] Logs are saved into C:\tmp\lmstudio-server-log.txt
[2023-12-17 11:43:45.050] [INFO] [LM STUDIO SERVER] Setting context overflow policy to: Rolling Window
Server Model Settings
Reset to Default
Preset
Vicuna v1.5 16K
The pre-prompt in the preset is ignored in server mode.
Danger Zone
Discard Changes
Override Preset
GPU Acceleration
Show Help
GPU Offload
n_gpu_layers
10
Detected GPU type (right click for options)
Nvidia CUDA
Context Length
n_ctx
2048
mistralai_mistral-... supports up to 32768 tokens.
Advanced Configuration
Config for Server Configuration
New Preset
Export to JSON
Prompt Format
Model Initialization
Hardware Settings
Tools
Model Inspector
Context Overflow Policy
v0.2.9
Model Downloads
0 downloading • 2 completed
Start/stop the server
Logs showing startup and connection information