- *Cluster singleton*—An actor that has only one instance in the entire cluster and is relocated to another node in case of failure. More on this later in this chapter.
- *Cluster sharding*—Actors are automatically distributed to different nodes by a partition key, which is also used to communicate with the actors. This provides load balancing and resilience when the nodes leave the cluster. More about this in chapter 9.
- *Distributed data*—Distributed key-value objects are ultimately consistent and can be used by any actor in any node of the cluster. The values are conflict-free replicated data types.
- *Distributed publish-subscribe*—Actors can create, publish, and subscribe to topics.
- *Reliable delivery*—Instead of the default at-most-once delivery, this module offers at-least-once guarantees.
- *Multi-DC cluster*—You can deploy clusters in different data centers (DCs) to increase resilience.

In addition to this plethora of modules, several tools under the Akka Management umbrella are designed to interact with cluster nodes:

- *Core* (or *Management*)—Provides the basic functionality to interact with the cluster through HTTP, which forms the base for different modules in Akka Management.
- *Management Cluster HTTP*—Uses the core module to offer an endpoint to read and modify the state of the nodes in the cluster.
- *Cluster Bootstrap*—Automates the discovery of nodes in the network instead of using static configuration.