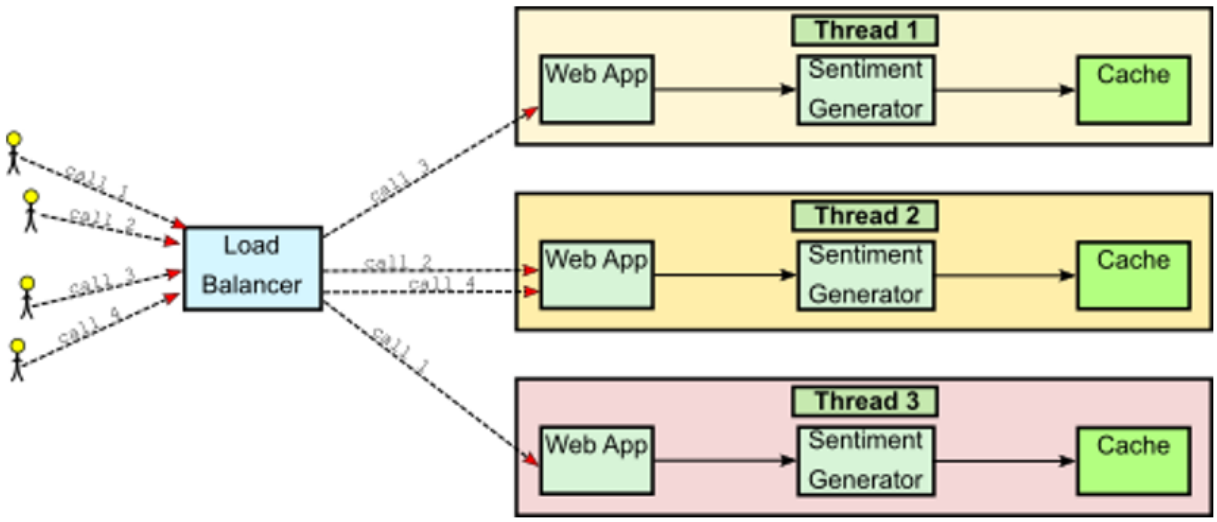


Figure 7.4 A possible configuration for the app running with one cache per thread. Each thread hosts a full pipeline (web server, sentiment generator and cache) and handles one call at the time, running synchronously. A load balancer makes sure to distribute the calls to the many threads running, waiting for each web server to answer before forwarding the next call (in figure, the web app in thread 2 would have returned the response for call #2 before the load balancer could forward call #4).

What happens, though, if in the example sequence above the second request for Google is processed while the first one is still computing the intermediate results?

If the first call hasn’t finished yet, no value has been added to the cache for Google. So, when the second call comes, the web app sees a cache miss, and calls the sentiment generator to retrieve all the data from social networks and compute the sentiment again.