Now that we have described how k-means is implemented, let’s take another look at how it works.

12.2.1 Issues With K-means

Figure 12.5 shows the result of applying k-means to an artificial dataset; the dataset has been carefully crafted to represent the ideal situation where k-means excels: the clusters are easily (and linearly[6]) separable, and they all are approximately spherical.

