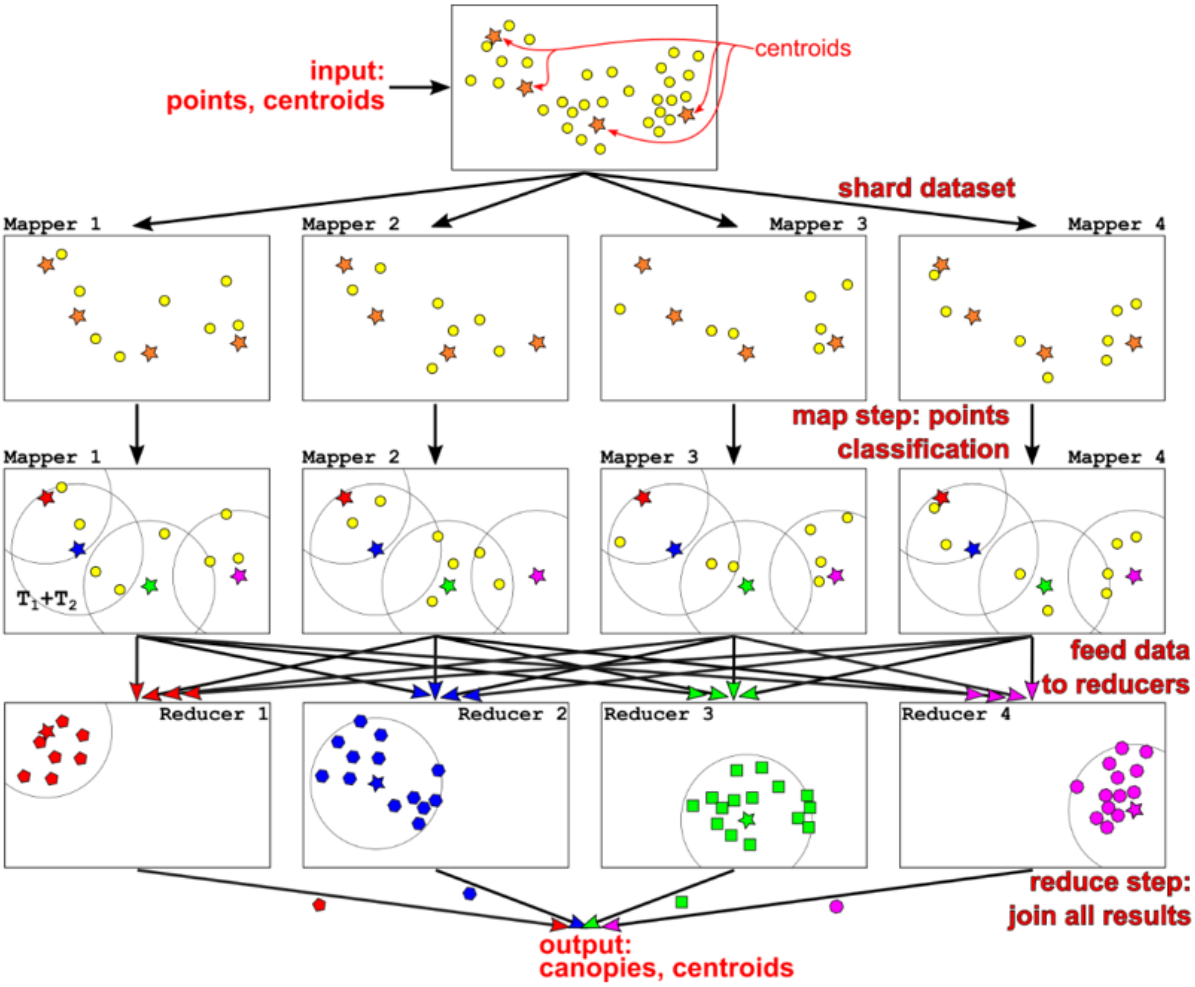


Figure 13.11 An example of a MapReduce job to classify points into canopies, given a list of centroids.

The method performing these assignments, `classifyPoints`[15], is run in the mappers for each shard separately: as for k-means, in fact, this step can be performed independently on each point, as long as the mapper has the full list of centroids.

The output of each mapper will be a list of lists with one entry per centroid: the list of points associated with that centroid; notice that each point can be associated with at least one, but potentially many centroids. Each mapper will