discard progressively more and more attempted updates, we'll need several iterations ("technically": a lot of them!) and, of course, a good function for the random steps to probe the problem space.

This makes the algorithm particularly slow and resource-consuming, especially if iterations (generating a random point and evaluating the cost function) are expensive to compute.

### 17.1.4 Short-range vs Long-range Transitions

Now, again, the question arises: do we want to progressively limit the range of transitions, and in turn how far the algorithm can move, at each step, in the domain space?

We could make the length of the maximum update step dependent on the temperature parameter `T`, and as such it would shrink in time.

But the interesting part is that even if we keep the same maximum step length for the whole process, there is some kind of indirect reduction while temperature cools: as shown in figure 17.6, the filtering actually happens on the codomain of the cost function, but the indirect effect is limiting the domain to a subset of the original problem space: and, in turn, the effect is the same as water running down funnels: it gets channeled, and velocity slows down (as pressure build up).