```
              \"database\": \"etherpad\",
              \"user\": \"etherpad\",
              \"password\": \"etherpad\"
            },
            \"exposeVersion\": true
          }
          " > settings.json                                          ⑧
          ./src/bin/run.sh &'                                        ⑨

          /opt/aws/bin/cfn-signal -e 0 --stack ${AWS::StackName}
➥ --resource AutoScalingGroup --region ${AWS::Region}               ⑩
# [...]
```

① The Auto Scaling group uses the launch template as a blueprint when launching EC2 instances.

② Picks the Amazon Linux 2 AMI dependent on the current region

③ Selects the instance type applicable for the Free Tier

④ Here is where the magic happens: the user data defined here is accessible by the EC2 instance during runtime.

⑤ The operating system will execute this bash script at the end of the boot process.

⑥ If any step fails, the script will abort and notify CloudFormation by calling cfn-signal.

⑦ Fetches Etherpad from the GitHub repository

⑧ Creates a settings file for Etherpad containing the database host name

⑨ Starts Etherpad

⑩ Notifies CloudFormation about a successful deployment

Let’s repeat how this deployment works:

1. The launch template specifies a script to add to the user data when launching an EC2 instance.
2. The Auto Scaling group launches an EC2 instance based on the Amazon Linux 2 image.
3. At the end of the boot process, the EC2 instance executes the bash script fetched from user data.
4. The bash script fetches, configures, and starts Etherpad.