availability goals are always satisfied and redistributing apps as nodes are added or removed from the cluster. An app is implemented by one or more *pods*, which encapsulate the software needed to run that app. A pod is itself made up of one or more Linux *containers*, each typically running a single process such as an HTTP API server.

> DEFINITION A Kubernetes ***node*** is a physical or virtual machine that forms part of the Kubernetes cluster. Each node runs one or more ***pods*** that implement apps running on the cluster. A pod is itself a collection of Linux ***containers*** and each container runs a single process such as an HTTP server.

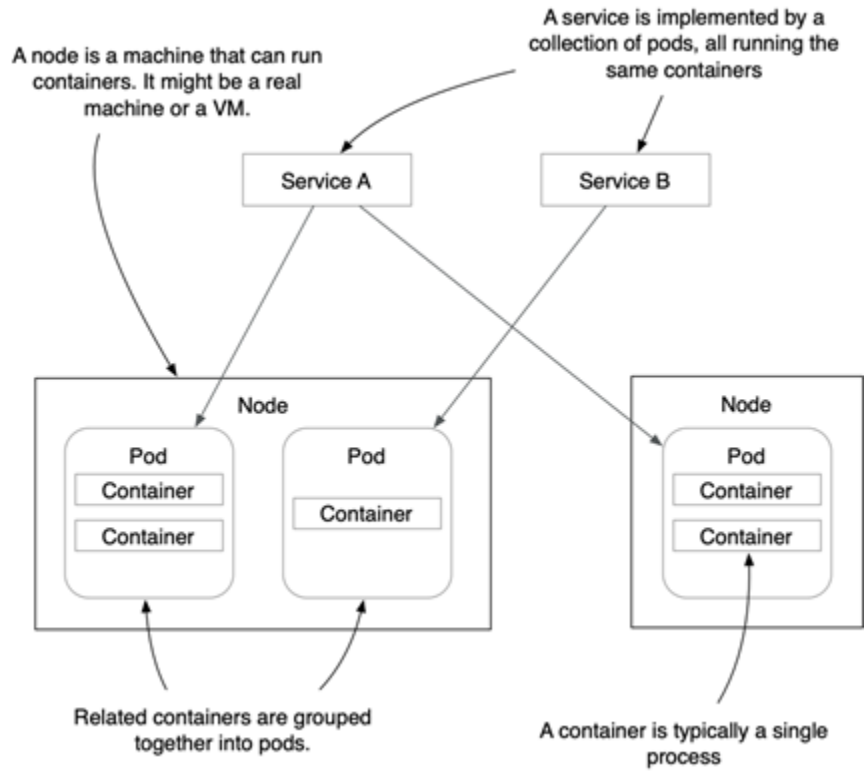


Figure 10.2 In Kubernetes an app is implemented by one or more identical pods running on physical or