# 17 AI in policy practice

## This chapter covers

- Why AI systems must not act as authorization decision-makers
- How AI assists policy authoring and design without becoming authoritative
- Using AI in runtime analysis to compare design-time intent with observed behavior
- Enforcing authorization before retrieval in RAG systems
- How architecture enables accountability and governance in AI-enabled systems

As AI systems become part of daily life, they play a bigger and bigger role in key decisions. These include what data is displayed, what actions are allowed, and who is accountable if something goes wrong. As a result, many teams are tempted to delegate authorization decisions to AI. That idea seems reasonable, but it’s usually a mistake.

Authorization exists to enforce the intent behind access decisions. It relies on determinism, explainability, auditability, and predictable failure modes. AI systems excel at many tasks, such as interpretation, synthesis, and exploration. When these strengths are applied appropriately, AI can significantly enhance the design, understanding, testing, and governance of policies without becoming the component responsible for access determination.

When ACME Corp started experimenting with AI tools, authorization was not the main focus. Teams around the company explored AI for various reasons: improving developer productivity, automating customer support, searching documents, analyzing data, and conducting initial tests with conversational interfaces. As these tools were integrated into existing systems, they began to interact with authorization in ways that were initially unclear.