reasoning about the effects of change. What it cannot do is determine which policy best aligns with organizational intent.

By keeping AI within the design-time workflow and outside the enforcement path, ACME was able to enhance policy quality without compromising the guarantees that authorization systems must ensure. Policies remained deterministic, explainable, and auditable, while engineers received quicker feedback and better tools for expressing intent.

17.3 What AI can tell you about your authorization policies

Writing a policy expresses intent; authorization systems enforce behavior. The distinction between the two often only becomes clear at runtime, when policies are evaluated against actual data, relationships, and context. As ACME increased its use of policy-based authorization across systems, this gap between design-time intent and runtime reality grew more evident.

Figure 17.3 illustrates this difference. On the left, policies are written to reflect organizational intent, often with AI assistance to draft, explain, or refine rules. On the right, authorization outcomes from production are recorded in decision logs and reviewed alongside the relevant policies. AI-assisted analysis helps teams compare what happened at runtime with what the policies were intended to express, surfacing gaps, surprises, and edge cases. Although AI is used in both scenarios, the inputs and objectives differ, as does the AI’s role. At design time, AI helps articulate intent. At runtime, it analyzes the results.