17.3.3 Asking better questions

As teams grew more confident with this approach, the type of questions they asked evolved. Instead of solely inquiring if a specific request was permitted, they started posing more general questions about policy behavior:

- “What is the effective audience for this document type?”
- “How often do exceptions override general rules?”
- “Which policies are responsible for most access decisions in this service?”

These questions were difficult to answer solely by reading policies. They required aggregating and interpreting runtime data. AI was well-suited for this because it could help people examine large amounts of information and find patterns worth exploring.

17.3.4 From observation to accountability

Understanding runtime behavior also played a key role in accountability. When access decisions were challenged, ACME could present concrete evidence: the policies in effect at the time, the data used during evaluation, and the decision that was made. AI-assisted analysis facilitated the collection, interpretation, and explanation of this evidence, particularly for non-technical stakeholders.

Authorization decisions still originated from the policy engine, based on policy and data. But AI assisted in explaining and analyzing those decisions afterward. This maintained the determinism and auditability needed for authorization while increasing the system’s transparency and clarity.

By using AI to bridge the gap between design-time intent and runtime reality, ACME gained a clearer understanding of how its authorization system actually functioned. That insight made policies easier to update, defend, and align with organizational goals. Next, we’ll see how these same principles apply when authorization is combined with retrieval-augmented generation systems, where enforcing access before retrieval becomes crucial.