17.4.4 Safe by construction, not safe by instruction

This architecture shifts responsibility for access control from prompts and model behavior to system design. Prompts still influence quality, tone, and task framing, but they no longer bear the responsibility of enforcing access boundaries. These boundaries are enforced structurally before any data retrieval.

For ACME, this distinction was crucial. Requiring authorization before retrieval prevented numerous failure modes and made governance easier. When questions arose about what an AI system could see or explain, the answer always traced back to policy and authorized scope, not prompt wording or model behavior.

By embedding authorization into the retrieval path, ACME made its RAG systems safe by construction rather than relying on best-effort controls. As AI systems advance beyond retrieval to autonomous action and delegation, this same principle becomes even more vital, creating a foundation for accountability in increasingly agent-driven architectures. We'll explore this more in the next chapter.

17.5 From architecture to accountability: how AI helps policy become practice

Earlier in this chapter, we discussed architectural boundaries: where AI belongs, where it doesn't, and how authorization limits what AI systems can access and do. Those boundaries are necessary but not sufficient. Authorization only gains meaning when decisions can be understood, explained, and acted upon over time. That is the purpose of accountability.

Figure 17.3 visually introduced this concept. Figure 17.5 reproduces the right side of that figure.