explains how observed authorization behavior becomes evidence, how that evidence guides review and oversight, and how AI helps close the loop between architectural intent and organizational responsibility.

At ACME, the shift to policy-based authorization laid a foundation for consistent decision-making, but it didn't automatically bring clarity. As policies increased in number and complexity, teams faced practical questions that were hard to answer by inspection: Who can access this resource today? Why was one request granted while another was denied? Which policies dominate outcomes in practice?

17.5.1 Accountability starts with observable decisions

Accountability goes beyond having well-written policies. It also requires evidence. As we discussed earlier, that evidence comes from authorization decisions recorded at runtime, linked to the policies and data that generated them. These decision logs provided a factual record of what the system allowed and denied, regardless of what anyone believed the policies were meant to do.

AI-assisted analysis made this record useful. By summarizing patterns in decision logs, highlighting outliers, and linking outcomes to specific policies, AI tools help teams turn anecdotes into evidence-based observations. Questions like “Who can actually access this resource?” or “Why did this decision happen?” can be examined directly through observed behavior.

This is the initial step in the governance process depicted in figure 17.5. The aim is not to automate judgment but to uncover information that would otherwise be hidden in logs.

17.5.2 Turning reports into responsibility

Once authorization outcomes become visible, responsibility is easier to identify. When authorization logic is embedded in application code or spread across multiple systems, it’s often unclear who is responsible for a specific decision. In contrast, policy-based authorization produces decisions that can be traced to specific policies and, consequently, to their owners.