AI-assisted analysis helps make these connections explicit. It can identify which policies are most frequently used, which ones lead to unexpected outcomes, and where exceptions override general rules. These reports do not make decisions, but they shape the questions governance bodies need to answer.

At ACME, this shifted governance conversations away from abstract discussions and toward concrete behavior. Instead of questioning whether a policy appeared reasonable, reviewers could assess if actual outcomes still aligned with organizational goals. Responsibility shifted from interpretation to accountability.

17.5.3 Closing the feedback loop

Figure 17.3 shows two feedback loops. The design-time loop refines policy intent through authoring and review. The runtime loop uses observed outcomes to inform governance processes.

Here we're focusing on the runtime loop, shown in figure 17.5, where AI supports analysis and reporting while humans make governance decisions. It helps summarize what happened, highlight what is unusual, and connect outcomes back to the policies that produced them. What it does not do is decide whether those outcomes are acceptable. That decision remains with humans in ACME's policy governance process, guided by evidence rather than intuition.

When governance determines that a policy no longer reflects desired behavior, changes are made intentionally and re-enter the design-time cycle. In this way, architecture, analysis, and governance support each other without collapsing into automated enforcement.

17.5.4 Why architecture makes accountability possible

None of this works without the architectural constraints outlined earlier in the chapter. If AI systems are allowed to infer access from prompts, retrieve data without permission, or act independently without policy review, the reports supporting governance lose their significance. Decisions become hard to explain and even harder to defend.