18 Authorization for AI agents

This chapter covers

- Why authorization becomes the hardest problem in agentic AI systems
- How AI agents interact with enterprise systems and their authorization requirements
- Policy-aware agent loops and runtime authorization checks
- Using policies to guide agent behavior, not simply to deny actions
- Delegation and subagent authority in agentic systems

A year after launching the Customer Collaboration platform, ACME started experimenting with a new class of tools powered by artificial intelligence (AI). What began as a few internal prototypes quickly spread throughout the organization. Engineers used agents to analyze logs and troubleshoot deployments. Sales teams utilized them to prepare account briefings. Customer support tested agents that could summarize cases and suggest responses.

The early results were impressive. Tasks that once took hours of manual effort could now be done in minutes. The engineering team developed integrations that let agents access internal systems, retrieve documents, and carry out routine operational tasks. The company started to envision an AI-enabled workforce where software agents worked alongside employees.

To coordinate these efforts, ACME hired Trudy, a seasoned technology leader, as Senior Director for AI. Trudy's mandate was simple in principle but ambitious in scope: help ACME deploy AI agents safely and effectively across the enterprise.

In her first weeks, Trudy met with teams across the company to understand how they were using these new tools. The conversations were lively. Teams described agents that could read documents, analyze customer data, prepare reports, and even trigger workflows.