GOVERNING AI AGENTS WITH AMAZON BEDROCK'S AGENTCORE POLICY

Amazon's introduction of AgentCore Policy in Amazon Bedrock (https://docs.aws.amazon.com/bedrock-agentcore/) shows a growing recognition that agent behavior must be governed using the same policy-based authorization patterns that have long been applied to human and service identities.

AgentCore Policy uses the Cedar policy language to specify what tools an agent can invoke and under which conditions. These policies are assessed outside the agent's runtime, usually at an enforcement gateway that controls access to external systems. This design enables organizations to externalize decision logic, maintain centralized governance, and ensure uniform enforcement across multiple agents and applications.

For example, an enterprise might allow a customer-support agent to retrieve account information but prohibit it from starting financial transactions. Another policy might permit a procurement agent to submit purchase requests only within specified spending limits. Because these policies are expressed declaratively and evaluated at runtime, they can adapt flexibly to changing contexts such as user identity, delegation relationships, or operational constraints.

AgentCore Policy marks a significant shift in how AI systems are built. Instead of embedding behavioral rules directly into agent code or prompts, organizations are now starting to treat agent authorization as a primary architectural concern. This change is in line with the larger move toward policy-based access control for dynamic systems, where decisions are handled externally and assessed on an ongoing basis.

As agent ecosystems become more complex—with agents calling on other agents, acting for humans, and engaging with sensitive enterprise systems—formal authorization methods like Cedar-based policy evaluation will play an increasingly important role in ensuring trust, safety, and operational accuracy.