The engineers responded by implementing strict policy controls on any operation capable of modifying the system's configuration. Tools that created or updated workflows, changed infrastructure settings, or altered authorization policies were restricted to only authorized principals. Agents acting under delegated authority were permitted to perform operational tasks, but they were explicitly forbidden from modifying the control plane.

For example, ACME's internal automation platform allows users and services to trigger predefined workflows. Agents can start existing workflows as part of their tasks but cannot create new workflows or modify existing ones. A Cedar policy enforcing this rule might look like the following:

```
forbid (
    principal in Agent,
    action in [
        Action::"automation.create",
        Action::"automation.update",
        Action::"automation.delete"
    ],
    resource in Entity::Automation
);
```

This policy prevents any agent principal from modifying automation definitions while still allowing agents to trigger approved workflows.

This approach effectively childproofed the system. Agents could interact with enterprise systems and carry out useful work, but the mechanisms that defined policy and governance remained protected.

As ACME gained more experience with agent behavior and the policies that govern it, the engineers were able to refine these restrictions. Some operations that were initially blocked could be safely delegated once the team understood the risks and implemented suitable safeguards. The core control-plane protections remained in place, but the operational boundaries around agents gradually expanded as the policy system and governance processes matured.

In practice, this meant defining two categories of actions: