- **operational actions**, which agents might perform as part of their tasks
- **control-plane actions**, which were reserved for trusted administrative principals

Policies enforced this separation, ensuring that even a powerful agent operating under delegated authority could not alter the rules that governed its behavior.

18.6 Delegation and subagents

As ACME's engineers grew more comfortable with their agent architecture, they began exploring more complex tasks that involved multiple agents working together. A common pattern quickly emerged: a primary agent would break a task into smaller pieces and assign them to specialized subagents.

For example, when an employee asks the ACME Assist agent to prepare a customer briefing, the agent might break the work into distinct steps. One subagent retrieves relevant documents, another analyzes recent customer activity, and a third combines and summarizes the results. Each subagent focuses on a specific part of the task.

This pattern raises an important authorization question: under whose authority are these subagents operating?

In traditional systems, delegation is often implicit. A service calls another service, and trust is assumed based on identity or configuration. In agent systems, that approach is insufficient. Because agents plan and act dynamically, implicit delegation risks expanding authority beyond what was intended. ACME addressed this by making delegation explicit and treating it as structured data that policies can evaluate.