Step 4: Evaluate options as a human reviewer

Review the proposed alternatives with an eye toward clarity, scope, and long-term maintainability. Ask questions such as:

- Which option communicates intent most clearly?
- Which minimizes the risk of unintended access?
- Which will be easiest to understand and modify later?

Although the alternatives may be functionally equivalent, they are not equivalent from a design perspective. Choosing between them is a human judgment, not a mechanical decision.

Step 5: Apply the change deliberately

Select one approach and apply it manually to the policy set. Update or add test cases to reflect the new behavior and re-evaluate the original request to confirm that the outcome now matches the stated intent.

If the result is surprising, revisit the policy change before proceeding. Unexpected outcomes are a signal to slow down, not to add more conditions.

Step 6: Reflect on the change

Finally, consider whether the modified policy still reads as a coherent statement of authority. A good policy change should be explainable in plain language and defensible in an audit.

This structured approach keeps AI assistance focused on exploration and drafting while ensuring that responsibility for authorization behavior remains firmly with the human reviewer.

Things to try

Once you are comfortable with this process, try applying it to a few small variations. Each of these exercises is intentionally narrow and can be