reversed easily.

1. **Introduce a temporary exception.** Modify the policy set so that the request is denied only under a specific condition, such as a particular document attribute or time window. Ask the AI to propose multiple ways to express the exception and compare their clarity and scope.
2. **Generalize the exception.** Instead of denying access for a single principal, explore how the same approach would apply to a small group. Use the AI to surface where additional relationships or attributes would be required, and decide whether that added complexity is justified.
3. **Revert and explain.** Restore the original policy behavior and ask the AI to explain the difference between the original and modified versions in plain language. This is a useful way to validate that the change was truly minimal and that its intent is easy to articulate.

These exercises reinforce the habit of starting from intent, exploring alternatives, and making deliberate, reviewable changes, rather than accumulating policy complexity through trial and error.

B.7 Guardrails and review checklist

The starter prompt that we loaded as part of getting set up is the foundation for how the AI interprets context, constraints, and its own role in the policy authoring process. It shapes not only what the AI can see, but how it reasons about policy, authority, and responsibility. Because of that, the prompt deserves careful attention—it encodes guardrails just as deliberately as the policies themselves.

The starter prompt used in this appendix includes several intentional constraints:

- **Use the schema and existing policies as the source of truth.** This ensures that AI suggestions are grounded in the actual authorization model. Any policy proposal that invents entities, actions, or attributes is incorrect by definition.