- **Do not invent new entities, actions, or attributes.** Schema changes are architectural decisions. Keeping them out of AI-assisted policy drafting prevents accidental expansion of the authorization model.
- **Prefer least-privilege outcomes by default.** This biases AI suggestions toward narrower grants, reducing the risk of unintended access.
- **Provide example requests and expected outcomes**. Requiring examples forces proposed changes to be evaluated in concrete terms rather than accepted as abstract logic.
- **AI output is advisory and will be reviewed by a human.** This explicitly reinforces accountability. The AI is a drafting assistant, not a decision-maker.

Taken together, these constraints shape the AI's behavior toward exploration and explanation rather than authority. They are not optional, and they should not be removed lightly.

The prompt establishes how the AI operates within defined boundaries. Humans use the following checklist to evaluate and accept or reject the output.

Confirm intent

- Can you explain the policy change in one or two plain-language sentences?
- Does the change do exactly what was intended, and nothing more?
- Would another reviewer reach the same conclusion after reading the policy?

Check scope

- Is access limited to the smallest set of principals, actions, and resources?
- Are there any new conditions or relationships that could grant unintended access?
- Does the change interact with other policies in unexpected ways?