Review structure

- Is the choice between a new forbid policy, a narrower when clause, or an unless clause deliberate?
- Does the policy read clearly, or is intent obscured by complexity?
- Will this policy still make sense to someone who was not involved in the change?

Validate behavior

- Are there test cases that demonstrate both intended permit and deny outcomes?
- Do the tests reflect intent, not just syntax?
- Does the evaluated result match what you would expect to explain during an audit?

Reaffirm accountability

- AI suggestions are advisory, not authoritative.
- The policy engine evaluates decisions, but humans remain accountable for what those decisions allow or deny.
- If you cannot confidently defend the change, do not accept it.

Starting with the prompt and ending with human accountability reinforces the central lesson of this appendix: guardrails are not something you add after using AI—they are what make AI assistance safe to use in the first place.

B.8 Key takeaways

This appendix has shown a practical, step-by-step approach to using AI tools to assist with authorization policy authoring. The focus was not on automating decisions, but on improving the quality and clarity of human-authored policies.

The key lessons are straightforward: