The agent does not simply stop when denied. It replans and adapts to the guardrails that policy sets.

C.3.5 What this demo shows

This demo illustrates the reactive authorization model:

- every tool invocation is checked before execution
- policy decisions are enforced consistently
- denied actions are fed back into the agent’s reasoning

This establishes a critical boundary: the agent can plan freely but can act only within policy-defined limits. At the same time, the agent remains flexible. It can adapt its behavior when actions are denied, allowing it to continue working toward a goal within those limits.

This model is simple and effective, but it has a limitation: the agent learns boundaries only after encountering them. If the agent repeatedly proposes disallowed actions, it will receive repeated denials. This limitation motivates the next demo.

C.4 Demo 2: constraint-aware planning

The first demonstration showed how authorization decisions are enforced at runtime. In that model, the agent proposes actions and learns from permit or deny responses. This demonstration extends that model by enabling the agent to consult the policy system during planning.

Instead of discovering constraints through failed actions, the agent can query the policy system to determine what is allowed before acting. This aligns with the constraint-aware planning pattern described in section 18.4.