scalability, and performance. In this case, you can take advantage of the many data services explicitly built for the cloud, such as Amazon DynamoDB, Azure Cosmos DB, and Google BigQuery.

Cloud native data services are a fascinating topic, but we will mainly deal with applications in this book. Data-related concerns like clustering, replication, consistency, and distributed transactions will not be detailed too much in the book. I would love to, but they deserve their own book to be adequately covered.

INTERACTIONS

Cloud native services communicate with each other to fulfill the requirements of the system. How that communication happens will affect the overall properties of the system. For example, choosing a request/response pattern (synchronous HTTP call) over an event-based approach (messages streamed through RabbitMQ) will result in different levels of resilience for the application. In this book, we'll use different types of interactions, learn the differences between them, and see when to use each approach.

Summary

- Cloud native applications are highly distributed systems that are specifically designed for and live in the cloud.
- The cloud is an IT infrastructure provided as a commodity in terms of computing, storage, and networking resources.
- In the cloud, users pay only for the actual resources they use.
- Cloud platforms deliver their services at different levels of abstraction: infrastructure (IaaS), container (CaaS), platform (PaaS), functions (FaaS), or software (SaaS).