❷ The type of object to create

❸ The name of the Ingress

❹ Configures the ingress controller responsible for managing this object

❺ Ingress rules for HTTP traffic

❻ A default rule for all requests

❼ The name of the Service object where traffic should be forwarded

❽ The port number for the Service where traffic should be forwarded

At this point we are ready to deploy Edge Service and the Ingress to the local Kubernetes cluster. Open a Terminal window, navigate to the Edge Service root folder (edge-service), and run the following command:

```
$ kubectl apply -f k8s
```

Let’s verify that the Ingress object has been created correctly with the following command:

```
$ kubectl get ingress

NAME               CLASS   HOSTS   PORTS   AGE
polar-ingress      nginx   *       80      21s
```

It’s time to test that Edge Service is correctly available through the Ingress. If you’re on Linux, you don’t need any further preparation steps. If you’re on macOS or Windows, open a new Terminal window and run the following command to expose your minikube cluster to your localhost. The command will continue running for the tunnel to be accessible, so make sure you keep the Terminal window open. The first time you run this command, you might be asked to input your machine’s password to authorize the tunneling to the cluster:

```
$ minikube tunnel --profile polar
```