Usually, systems define SLOs and SLIs, such as with Kubernetes. See "Kubernetes scalability and performance SLIs/SLOs" (http://mng.bz/MBEm) for an example. Last but not least, if you want to start contributing in the space or just learn more, consider having a look at the excellent SLOconf (https://www.nobl9.com/community/sloconf) event, an annual gathering of the SRE community to exchange SLO-related updates and learnings.

## Summary

- Reliability regulations are on the horizon, like security was 15 years ago.
- Using SLOs to measure and report reliability for compliance is becoming a hot topic in finance, communications, and aeronautics and will likely spread to other industries.
- To make contractual promises in SLAs quantifiable and automated, define SLIs to drive the goals (SLOs).
- SLAs are contractual constructs between a service provider and a service consumer.
- SLOs are goals your organization defines and can be used to inform SLAs.
- SLIs are the actual measurement or metric (depending on the service type) to verify to what extent a goal (the SLO) was fulfilled.
- There is open source tooling (Pyrra and Sloth) available to implement SLOs/SLIs for the Prometheus ecosystem.
- There are some commercial offerings for SLOs, especially Nobl9's tool.
- You don't have to go all in with SLOs (e.g., error budgets or burn rate); in fact, not many folks in the wild do so (yet), but defining service metrics (SLIs) and tracking them against established goals (SLOs) helps you and your customer.
- SLOs can be used as a means of improving alert accuracy—early warning systems (nonemergency) that reduce false positives and the number of one-off rules used to suppress alerts and improve trend-based early warnings that can be tied to runbooks and automated remediation.
- Make sure to build on open specifications—OpenSLO, in this space.