```
exporters:
  otlp:
    endpoint: "jaeger:4317"
    tls:
      insecure: true
  logging:
  prometheus:
    endpoint: "otelcol:9464"
    resource_to_telemetry_conversion:
      enabled: true
    enable_open_metrics: true

processors:
  batch:
  transform:
    metric_statements:
    - context: metric
      statements:
      - set(description, "Measures the duration of inbound HTTP requests")
        where name == "http.server.duration"
  filter:
    metrics:
      exclude:
        match_type: strict
        metric_names:
          - queueSize

connectors:
  spanmetrics:

service:
  pipelines:
    traces:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlp, logging, spanmetrics]
    metrics:
      receivers: [otlp, spanmetrics]
      processors: [filter, transform, batch]
      exporters: [prometheus, logging]
    logs:
      receivers: [otlp]
      processors: [batch]
      exporters: [logging]
```

We have three exporters: Prometheus, Jaeger, and logging (to stdout).

We're using the spanmetrics connector to derive RED metrics from spans.

The pipelines are configured, with one for traces (connecting to metrics), one for metrics, and one for logs (to stdout only).

You can run the OpenTelemetry demo in several environments; we will be using the Kubernetes (http://mng.bz/Ao5Q) deployment in the following walk-through. But first, we need to make sure we have everything in place so that we can run the demo.